

# Chapter 17

Indexes

9 March 2017

Chris Sexton

# TOC

- Introduction
- Primary Indexes
- Culster Indexes
- Secondary Indexes
- Multilevel Indexes
- B+ Trees
- When to apply indexing

# Indexes

- primary mechanism for improving performance on a database
- takes advantage of the persistent data structures on disk
- many implementation details and decisions must be accounted for

# Utility

- Indexes avoid full table scans by allowing fast access to tuples
- This is a difference between  $O(n)$  and  $O(1)$  execution time, or...
- If you have 5,000,000 entries in a database with an access time of 1 ms per entry:

W|A: 5,000,000ms in minutes (<http://www.wolframalpha.com/input/?i=5,000,000+ms+in+minutes>)

- Versus instantaneous  $O(1)$  constant time lookup of the entry with a hash table
- Or  $O(\lg n)$  (Or,  $\sim 22$  operations) access with a tree or binary search on an ordered file

# Primary Index

- A **primary index** is an ordered file with records of fixed length of two fields
- Used for primary key access where we are looking up an exact record
- Use binary search on the fixed-length ordered file to find entries
- Notation:  $\langle K(i), P(i) \rangle$ : The first field,  $K(i)$  is the hashed key,  $P(i)$  is the pointer to the block containing the entry,  $i$
- This looks up entries by block, meaning the index has fewer entries than the actual data file
- Since the data file is ordered by this key, retrieving a block corresponding to the key value is sufficient

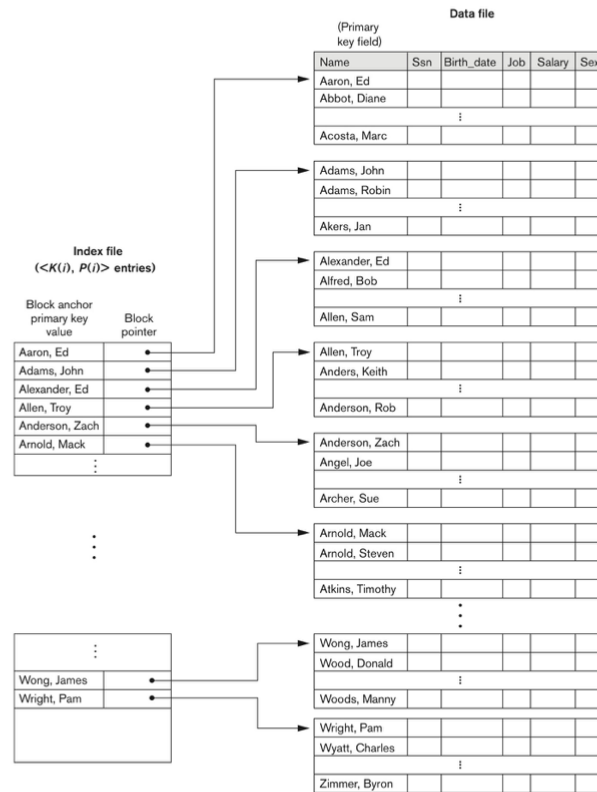
$\langle K(1) = (\text{Aaron}, \text{Ed}), P(1) = \text{address of block 1} \rangle$

$\langle K(2) = (\text{Adams}, \text{John}), P(1) = \text{address of block 2} \rangle$

$\langle K(3) = (\text{Alexander}, \text{Ed}), P(1) = \text{address of block 3} \rangle$

# Primary Index

**Figure 17.1** Primary index on the ordering key field of the file shown in Figure 16.7.



# Primary Index

## \* Gains

- The index file is strictly smaller than the data file
- Searching the index file is ordered, so fewer disk/block accesses are necessary

## \* Issues

- Insertion/Deletion is a problem in ordered files
- We must "mark" entries for deletion
- Insertion requires we move all entries "further" in the data file
- Can use overflow linked lists or blocks to alleviate these issues

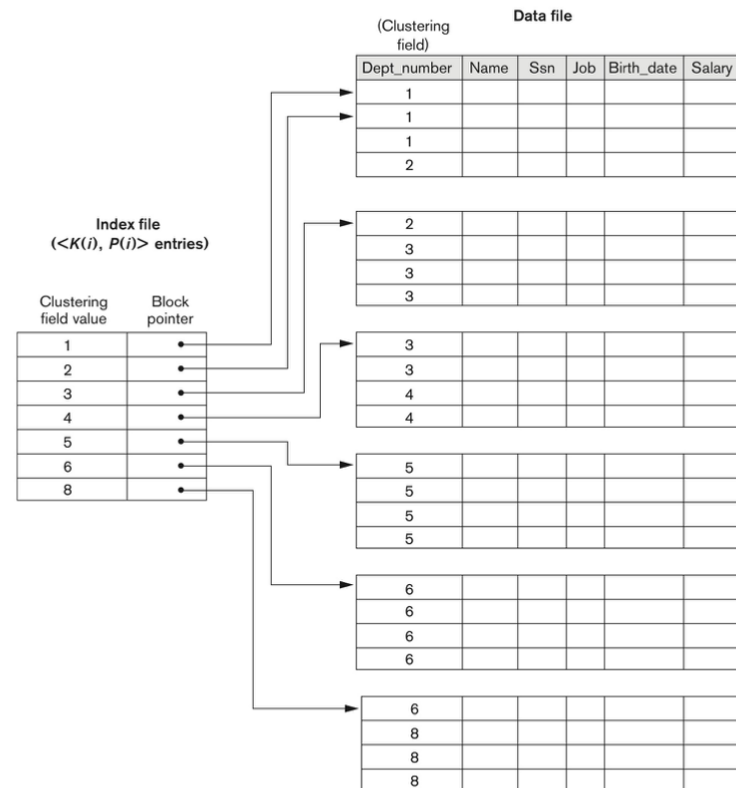
## Clustering Indexes

- Used for non-key/non-unique fields
- Speeds up retrieval of records that have the same value for a particular field
- Utilizes the same method for storage as the primary index, except the block pointer points to an ordered data file



# Clustering Indexes

**Figure 17.2** A clustering index on the Dept\_number ordering nonkey field of an EMPLOYEE file.



## Secondary Indexes

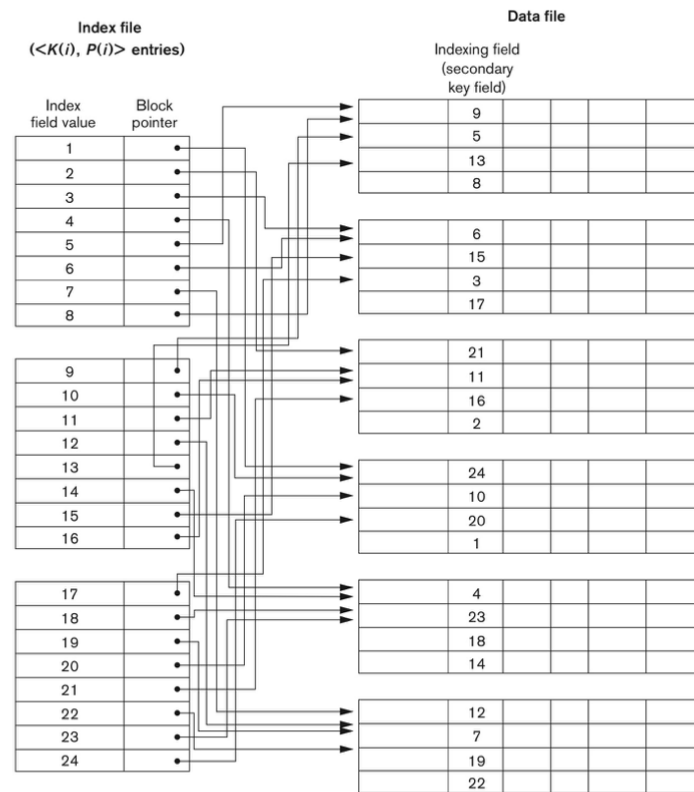
- Any field indexed that is not the primary means of accessing a particular set of data
- Could be ordered or unordered; unique or non-unique
- Can have many secondary indexes for any particular data set

## Unique Secondary Indexes

- Called a **secondary key**
- Corresponds to any field with the UNIQUE attribute set
- Cannot use a block-level key mechanism since the data file is not ordered by these fields
- Therefore, uses more storage space than the primary key
- We would have to do a linear search on the data without a secondary key, but still can perform binary search with the key

# Unique Secondary Indexes

**Figure 17.4** A dense secondary index (with block pointers) on a nonordering key field of a file.

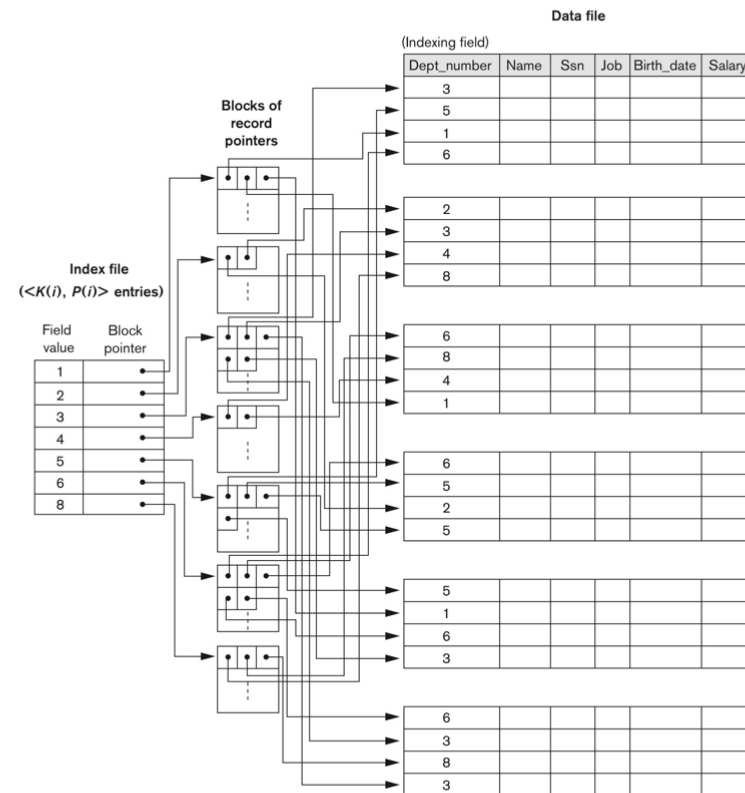


## Non-Unique Secondary Indexes

- Secondary index on an entry which is not a primary key
- Uses a table of lookup values matching the field  $K(i)$  pointing to a collection of blocks related to that value
- Extra level of indirection, meaning extra time cost

# Non-Unique Secondary Indexes

**Figure 17.5** A secondary index (with record pointers) on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values.



# Index Types

**Table 17.2** Properties of Index Types

Type of Index	Number of (First-Level) Index Entries	Dense or Nondense (Sparse)	Block Anchoring on the Data File
Primary	Number of blocks in data file	Nondense	Yes
Clustering	Number of distinct index field values	Nondense	Yes/no <sup>a</sup>
Secondary (key)	Number of records in data file	Dense	No
Secondary (nonkey)	Number of records <sup>b</sup> or number of distinct index field values <sup>c</sup>	Dense or Nondense	No

<sup>a</sup>Yes if every distinct value of the ordering field starts a new block; no otherwise.

<sup>b</sup>For option 1.

<sup>c</sup>For options 2 and 3.

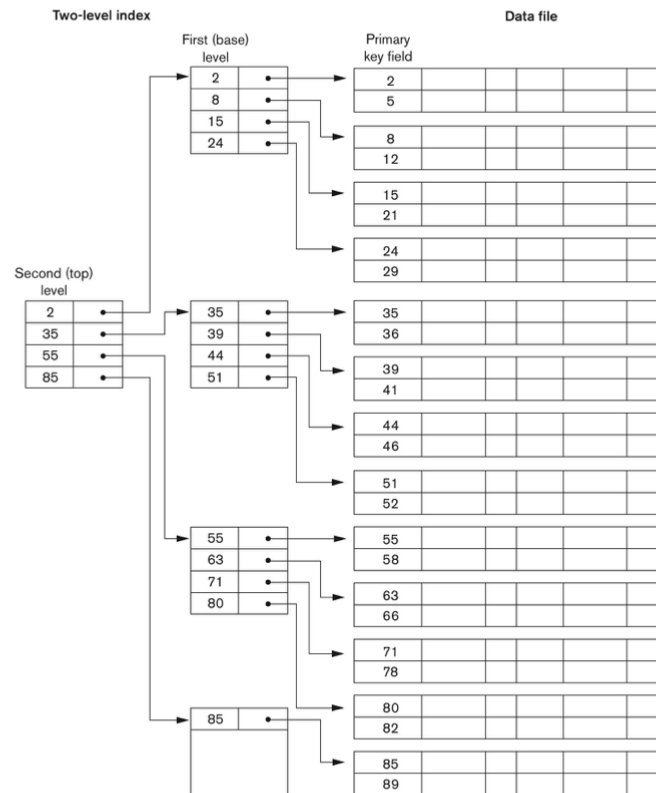
## Multilevel Indexes

- All previous indexes were on ordered index files with binary search
- A **multilivel index** reduces the search space even further by reducing disk reads
- Works for all previous types of indexes as long as the index has distinct values for  $K(i)$  and is of fixed-length



# Multilevel Indexes

**Figure 17.6** A two-level primary index resembling ISAM (indexed sequential access method) organization.



# B-Trees

- A search tree is used to guide search for records given a particular field value
- Different from a multilevel index, a search tree allows ordering and dynamic insertion/deletion
- B-trees are balanced trees that minimize wasted space

## \* Complexity

- Space required is  $O(n)$
- Insertion time:  $O(\lg n)$
- Search time:  $O(\lg n)$
- Removal time:  $O(\lg n)$
- Ranged query (e.g.  $1 < x < 5$ ) of  $k$  elements:  $O(\lg n + k)$

B+ Tree Basics (<https://www.youtube.com/watch?v=CYKRMz8yzVU>)

# Implementation

## \* Balanced Trees

- B Trees and B+ trees
- Allow searching for values within an index

## \* Hash tables

- key/value constant lookup
- Useful with exact lookups

## \* Other methods

- many variations depending on DBMS

## An example

```
select name  
from student  
where student_id = 38732
```

The student\_id field is a unique primary key and has an automatic index.

## An example

```
select name  
from student  
where name = 'George' and gpa > 3.5
```

- An index on the name may be a hash or tree
- An index on the GPA must be tree based to allow for searching values

## An example

```
select s.name, d.name  
from student s, department d  
where s.dept = d.id
```

- Query planning & optimization allows for selection of available indexes
- Index on student's department and department's primary key
- Database must have had the student's department specified as an index

## Downsides

- More space required for index storage
- Creation of index takes time
- Maintenance of index (on insert,update/delete) can take time and negate benefits

# When to index

## \* Tuning indexes

- Must verify indexes are used
- Queries not hitting an index may use full table scans and take a long time to run
- Indexes on tables that change frequently may slow the system down due to index updates/rebuilds

## \* Choosing indexes is an exercise in judgement.

- How big is the table?
- What is the distribution of data?
- Is the data mostly being queried or updated? Do many deletes occur?



# SQL Syntax for indexing

```
CREATE INDEX name ON table(attr)
```

```
CREATE INDEX name ON table(attr1, attr2, ... attrN)
```

```
CREATE UNIQUE INDEX name ON table(attr)
```

```
DROP INDEX name
```

- Used when creating a table
- Will start an indexing process for previously created tables (and incur load)

Thank you

Chris Sexton

